

Lecture 2: Agent Architectures

Multi-Agent Systems


Universitat Rovira i Virgili

Outline


1. Intelligent agent - definition
2. Kinds of environments
3. Agent architectures:
 1. Reactive
 2. Deliberative
 3. Hybrid

1. Intelligent agent - definition

- "An agent is anything that can be viewed as **perceiving its environment through sensors** and **acting upon that environment through effectors**"
- "Autonomous agents are computational systems that inhabit some **complex dynamic environment**, **sense** and **act** autonomously in this environment, and by doing so realize a set of **goals or tasks** for which they are designed"
- "An autonomous agent is a **system situated within and part of an environment that senses that environment and acts on it**, over time, in pursuit of **its own agenda** and so as to affect what it senses in the future"



Basic abstract view of an agent



Properties

- An agent has to be able to **react** (adapt its behaviour) in an appropriate way to the dynamic changes in its “**environment**”
 - Other computational agents
 - Human users
 - External information sources (e.g., sensors)
 - Physical objects (e.g., robots)
 - World Wide Web
 - ...
- This is one of several properties that an intelligent agent should have ... **[more on that next week]**

2. Kinds of environments

■ Accessible vs inaccessible

- An accessible environment is one in which the agent can obtain **complete, accurate, up-to-date information** about the environment's state
- Most moderately complex environments (including, for example, the everyday physical world and the Web) are inaccessible
- The more accessible an environment is, the simpler it is to build agents to operate in it



2. Kinds of environments

- Deterministic vs non-deterministic

- A deterministic environment is one in which **any action has a single guaranteed effect** — there is no uncertainty about the state that will result from performing an action
- The physical world can to all intents and purposes be regarded as non-deterministic
- Non-deterministic environments present greater problems for the agent designer

2. Kinds of environments

- **Episodic vs non-episodic**
 - In an episodic environment, the performance of an agent is dependent on a number of **discrete, independent episodes**, with no link between the performance of an agent in different scenarios
 - Episodic environments are simpler from the agent developer's perspective because the agent can decide which action to perform based only on the current episode — it does not need to reason about the interactions between different episodes

2. Kinds of environments

■ Static vs dynamic

- A static environment is one that can be assumed to **remain unchanged** except by the performance of actions by the agent
- A dynamic environment is one that has other processes operating on it, and which hence changes in ways beyond the agent's control
- The physical world and the Web are highly dynamic environments
- It is hard to design and implement agents in dynamic environments

2. Kinds of environments

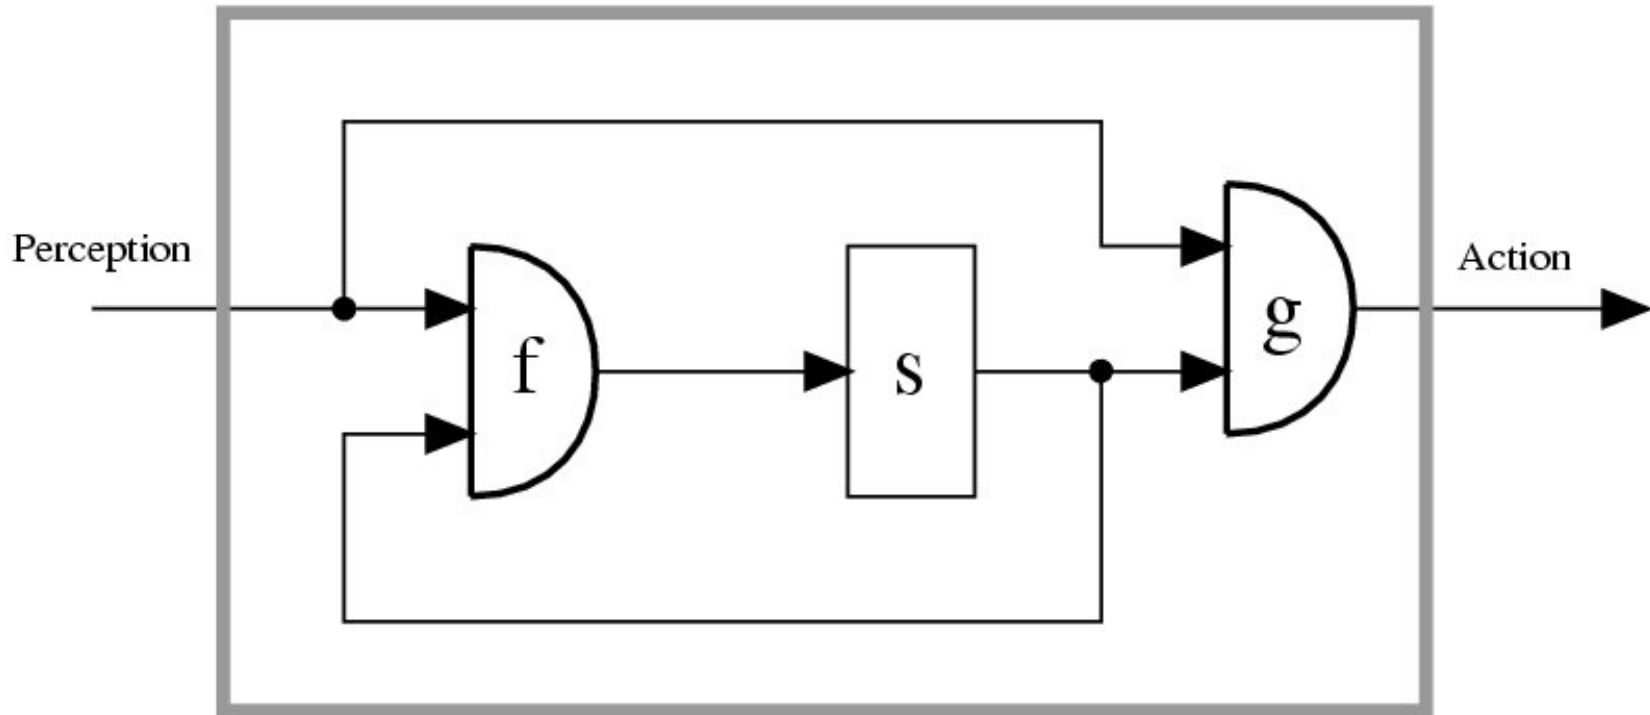
- Discrete vs continuous

- An environment is discrete if there is a **fixed, finite number of actions and percepts** in it
- The real world is a continuous environment
- Discrete environments are much simpler than continuous ones

3. Agent architectures

- An **architecture** proposes a particular **methodology** for building an autonomous agent
 - How the construction of the agent can be decomposed into the construction of a set of **component modules**
 - How these modules should be made to **interact**
- These two aspects define how the **sensor data** and the **current internal state** of the agent determine the **actions** (effector outputs) and the **future internal state** of the agent

From perception to action



f = state update function

s = internal state

g = output function

Main kinds of agent architectures

- **Reactive** architectures
 - Focused on fast reactions/responses to changes detected in the environment
- **Deliberative** architectures (symbolic)
 - Focused on long-term planning of actions, centred on a set of goals
- **Hybrid** architectures
 - Combining a reactive side and a deliberative side

Example: Reactive vs Deliberative

- Robot that has to reach a certain point
 - **Reactive**
 - Sensor in the front of the robot
 - Without any representation of the environment
 - Example: robot changes movement right/left when sensor detects obstacle or it reaches a crossroad. Minimal computation based on current location and destination point
 - **Deliberative**
 - Explicit representation of the environment (map)
 - Planning procedure that finds the minimal route between the current position and the destination
 - High computational cost
 - Optimal route (in static environments)
 - Possible dynamic re-plannings needed

3.1. Agent architectures: Reactive – basic ideas

- Reactive agents have
 - At most a **very simple internal representation** of the world
 - A **tight coupling of perception and action**
 - Behaviour-based paradigm
 - Intelligence is a product of the **interaction** between an agent and its environment

Classical example: ant colony

- A single ant has very little intelligence, computing power or reasoning abilities
- The union of a set of ants and the interaction between them allows the formation of a highly complex, structured and efficient system



Swarm intelligence



- Agents and virtual swarms
- Example: morphogenesis

Reactive agents – Main characteristics

- Emergent functionality
 - Simple agents
 - Simple interaction
 - Complex behaviour patterns appear as a result of the dynamic interactions
 - The global behaviour of the system is not specified a priori
 - Dynamic movement of robots, depending on obstacles

Reactive agents – Main characteristics

- Task decomposition
 - Agents composed of autonomous modules
 - Each module manages a given task
 - Sensor, control, computations
 - Minimal, low-level communication between modules
- There isn't any world global model
- There isn't any “planning/controller/coordinator agent”

Reactive agents – Main characteristics

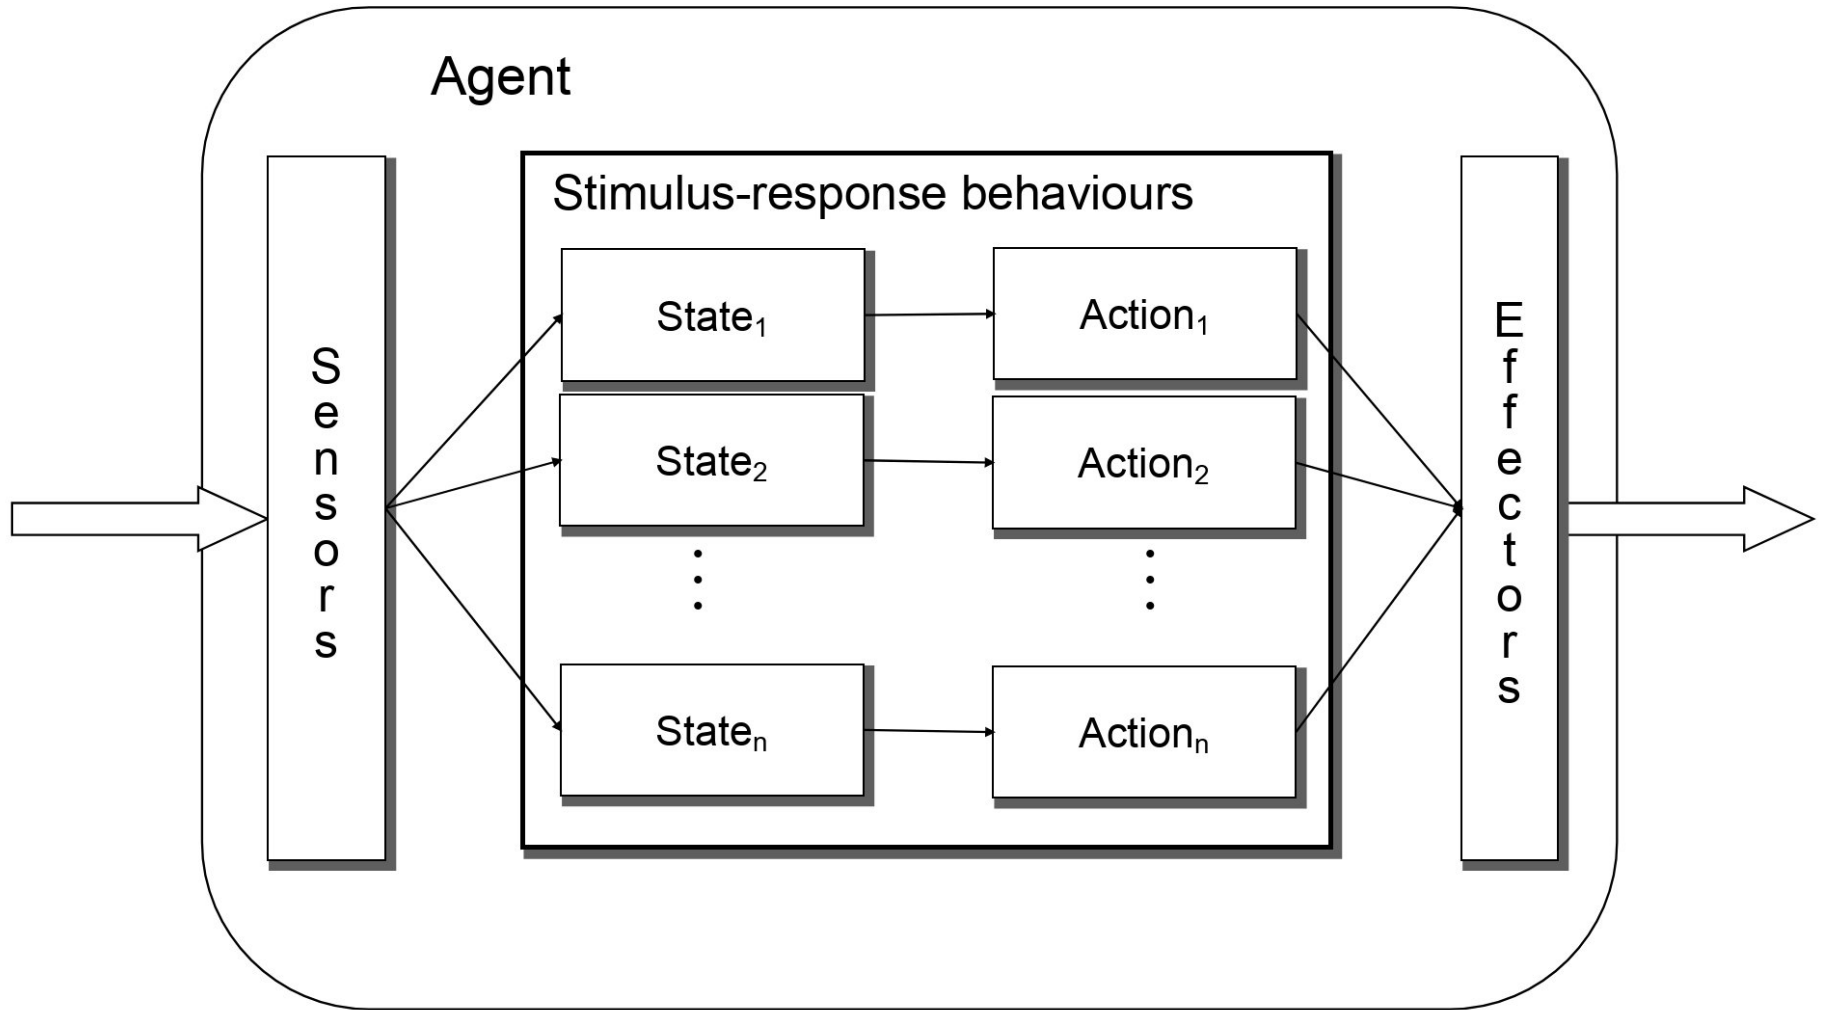
- Raw data
 - Basic data from sensors
 - There isn't any complex symbolic management of data as in classical AI
 - Refusal of the *Hypothesis of the physical symbols system* [basic pillar of symbolic AI]
 - “Intelligent behaviour can only be obtained in symbol-processing systems”

Reactive agents – Basic concept

- Each behaviour continually maps perceptual input to action output
- Reactive behaviour: **action rules: $S \rightarrow A$**
where S denotes the **states** of the environment, and A the primitive **actions** the agent is capable of performing
- Example:

$$action(s) = \begin{cases} Heater\ off, & \text{if temperature is over 15} \\ Heater\ on, & \text{otherwise} \end{cases}$$

Basic schema of reactive architecture



Rodney Brooks

Director of the Computer Science and Artificial Intelligence Lab (MIT) 1997-2007

- Roomba's i-robot father



Brooks refutal of symbolic AI

- Brooks put forward three theses:
 1. Intelligent behaviour can be generated *without explicit representations* of the kind that symbolic AI proposes
 2. Intelligent behaviour can be generated *without explicit abstract reasoning* of the kind that symbolic AI proposes
 - Reduced computation on sensor-like data
 3. Intelligence is an *emergent* property of certain complex systems. Intelligence is 'in the eye of the beholder'; it is not an innate, isolated property

Brooks – key ideas


- **Situatedness:** ‘Real’ intelligence is situated in the world
 - The world is its best model
 - The world is always up-to-date
 - A model is an abstraction, a simplification of the world, considering a particular set of characteristics and disregarding others
- **Embodiment:** ‘Real’ intelligence requires a physical body, and cannot be found in disembodied systems such as theorem provers or expert systems
 - These are called Physical Agents

Brooks – subsumption hierarchy

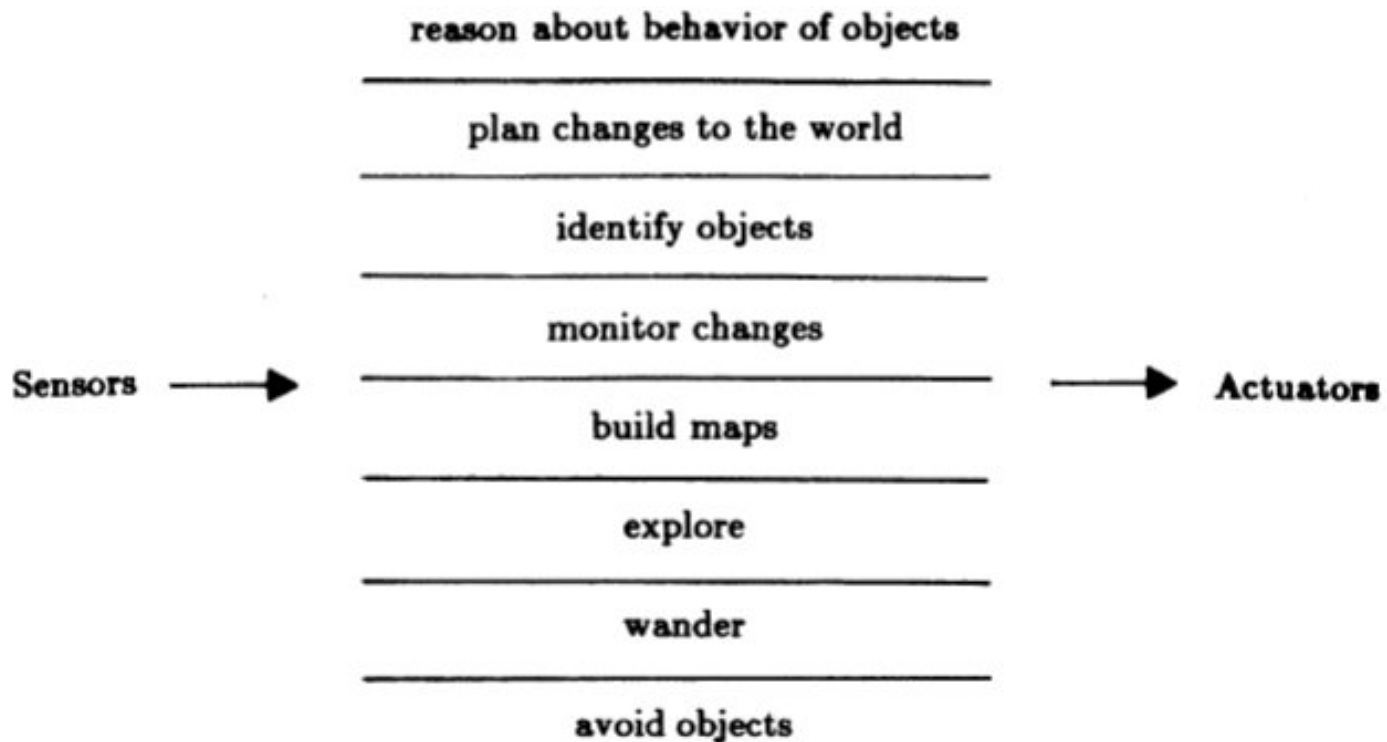
- To illustrate his ideas, Brooks built some systems based on his *subsumption architecture*
- A subsumption architecture is a hierarchy of task-accomplishing *behaviours*
- Each behaviour is a rather simple *rule-like structure*
- Each behaviour ‘competes’ with others to exercise control over the agent, as different behaviours may be applicable at the same time

Behaviour layers

- Lower layers represent primitive kinds of behaviour (such as avoiding obstacles)
- Higher layers represent more complex behaviours (e.g., identifying an object)
- Lower layers have precedence over layers further up the hierarchy
- The resulting systems are, in terms of the amount of computation they do, *extremely* simple
- Some of the robots do tasks that would be impressive if they were accomplished by symbolic AI systems

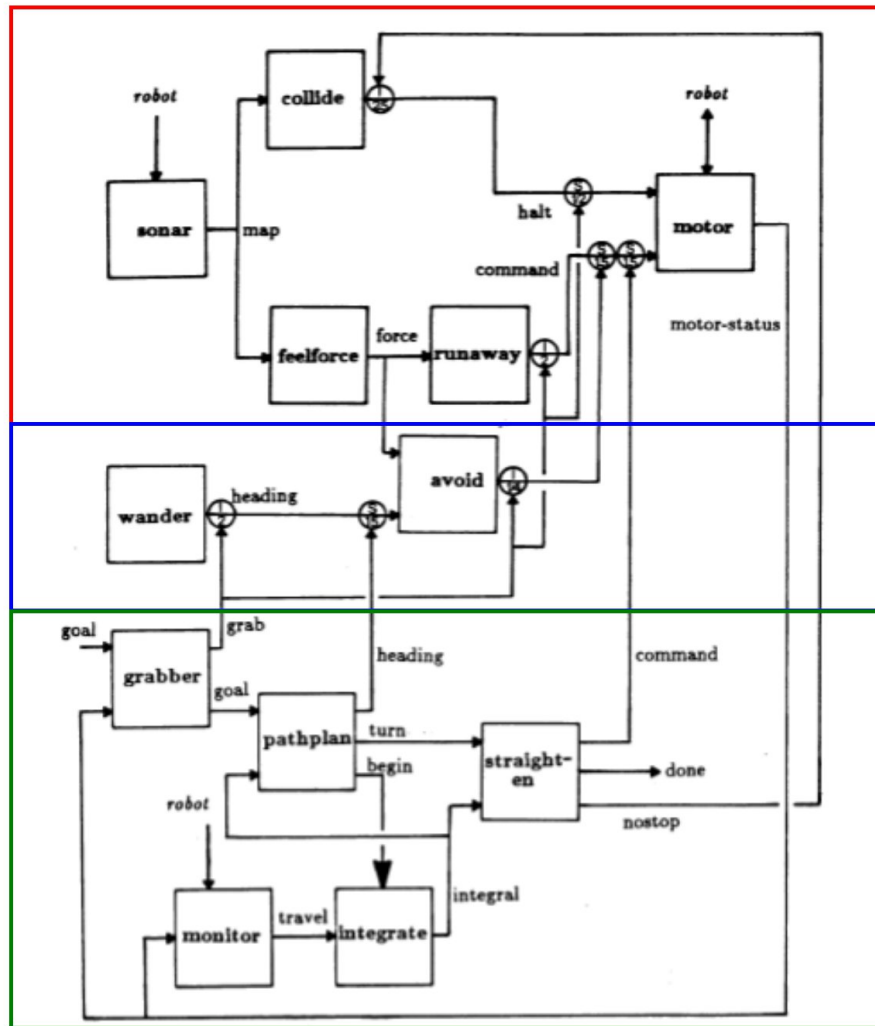


Decomposition Based on Task-Achieving Behaviours



From Brooks, "A Robust Layered Control System for a Mobile Robot", 1985

Levels 0, 1 and 2 Control Systems




Avoid collisions

Random movement

Pick up objects, monitor, explore

From Brooks, "A Robust Layered Control System for a Mobile Robot", 1985



Some robots from MIT Lab



Example: Luc Steels' Mars Explorer

- Steels' Mars explorer system, using the subsumption architecture, achieves near-optimal cooperative performance in simulated 'rock gathering on Mars' domain:

The objective is to explore a distant planet, and in particular, to collect sample of a precious rock. The location of the samples is not known in advance, but it is known that they tend to be clustered



Steels' Mars Explorer Rules

- For individual (non-cooperative) agents, the lowest-level behaviour, (and hence the behaviour with the highest “priority”) is obstacle avoidance:
if detect an obstacle then change direction (1)
- Any samples carried by agents are dropped back at the mother-ship:
*if carrying samples and at the base
then drop samples* (2)
- Agents will collect samples they find:
if detect a sample then pick sample up (3)

Steels' Mars Explorer Rules

- Agents carrying samples return to the mother-ship:
if carrying samples and not at the base
then move towards base and drop 2 crumbs (4)
- Agents can detect trails of crumbs:
if not(carrying samples) and detect crumb
then pick up 1 crumb and move away from base (5)
- An agent with “nothing better to do” will explore randomly:
if true then move randomly (6)

[Simulation](#)

Advantages of Reactive Agents

- **Simplicity** of individual agents
- **Flexibility, adaptability**
 - Ideal in very dynamic and unpredictable environments
- **Low computational cost**
 - Avoiding complex planning/reasoning procedures
 - Avoiding continuous model update
- **Robustness** against failure
 - No central planning component (e.g., ant colony)
- **Elegance**

Limitations of Reactive Agents

- Agents without environment models must have sufficient information available from the local environment
- If decisions are based on *local* environment, how can we take into account *non-local* information?
 - “Short-term” view
- No long-term planning capabilities
- **Limited applicability**
 - Games, simulations, basic robots (insects)

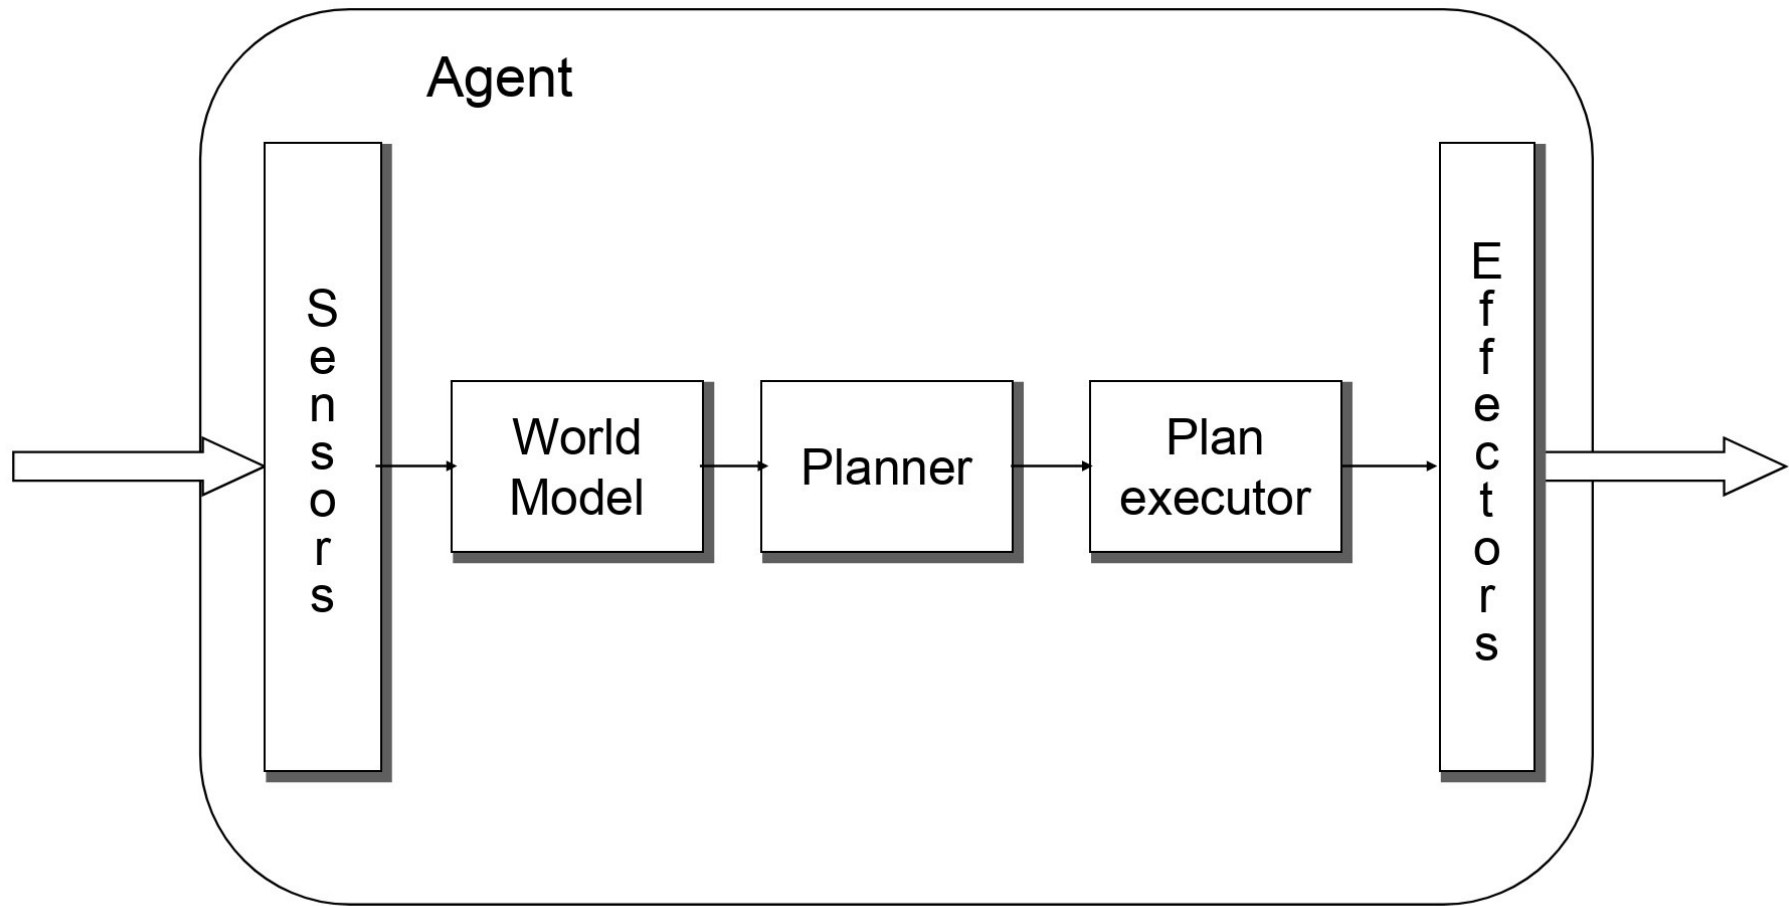
Limitations of Reactive Agents

- Difficult to make reactive agents that learn
 - Dynamic evolution of rules?
- Since behaviour emerges from component interactions plus environment, it is hard to see how to *engineer* specific agents (no principled methodology exists)
- It is hard to engineer agents with large numbers of behaviours as the dynamics of interactions become too complex to understand

3.2. Agent architectures: Deliberative


- **Explicit symbolic model** of the world
- Decisions are made via logical **reasoning**, based on pattern matching and symbolic manipulation
- **Sense-plan-act** problem-solving paradigm of classical AI planning systems

Basic schema of deliberative architecture



Belief-Desire-Intention (BDI) model

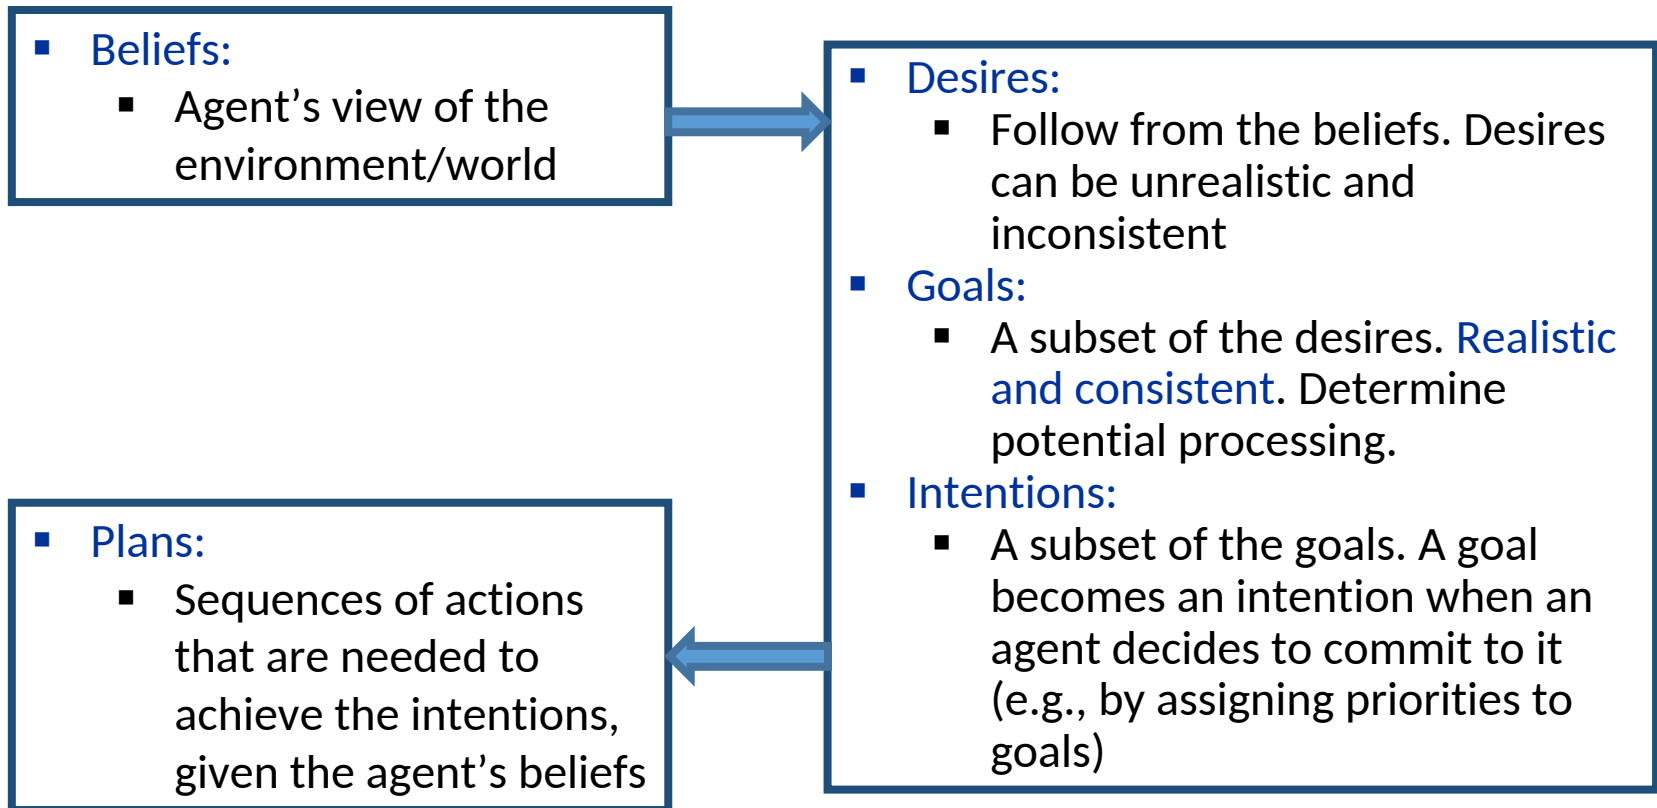
- A theory of practical reasoning
- Originally developed by Michael E. Bratman in his book "*Intentions, Plans, and Practical Reason*", (1987)
- Concentrates in the roles of the intentions in practical reasoning



Human practical reasoning

- Human practical reasoning consists of two activities:
 - *Deliberation*, deciding *which* state of affairs we want to achieve
 - the outputs of deliberation are *intentions*
 - *Means-ends reasoning*, deciding *how* to achieve these states of affairs
 - the outputs of means-ends reasoning are *plans*

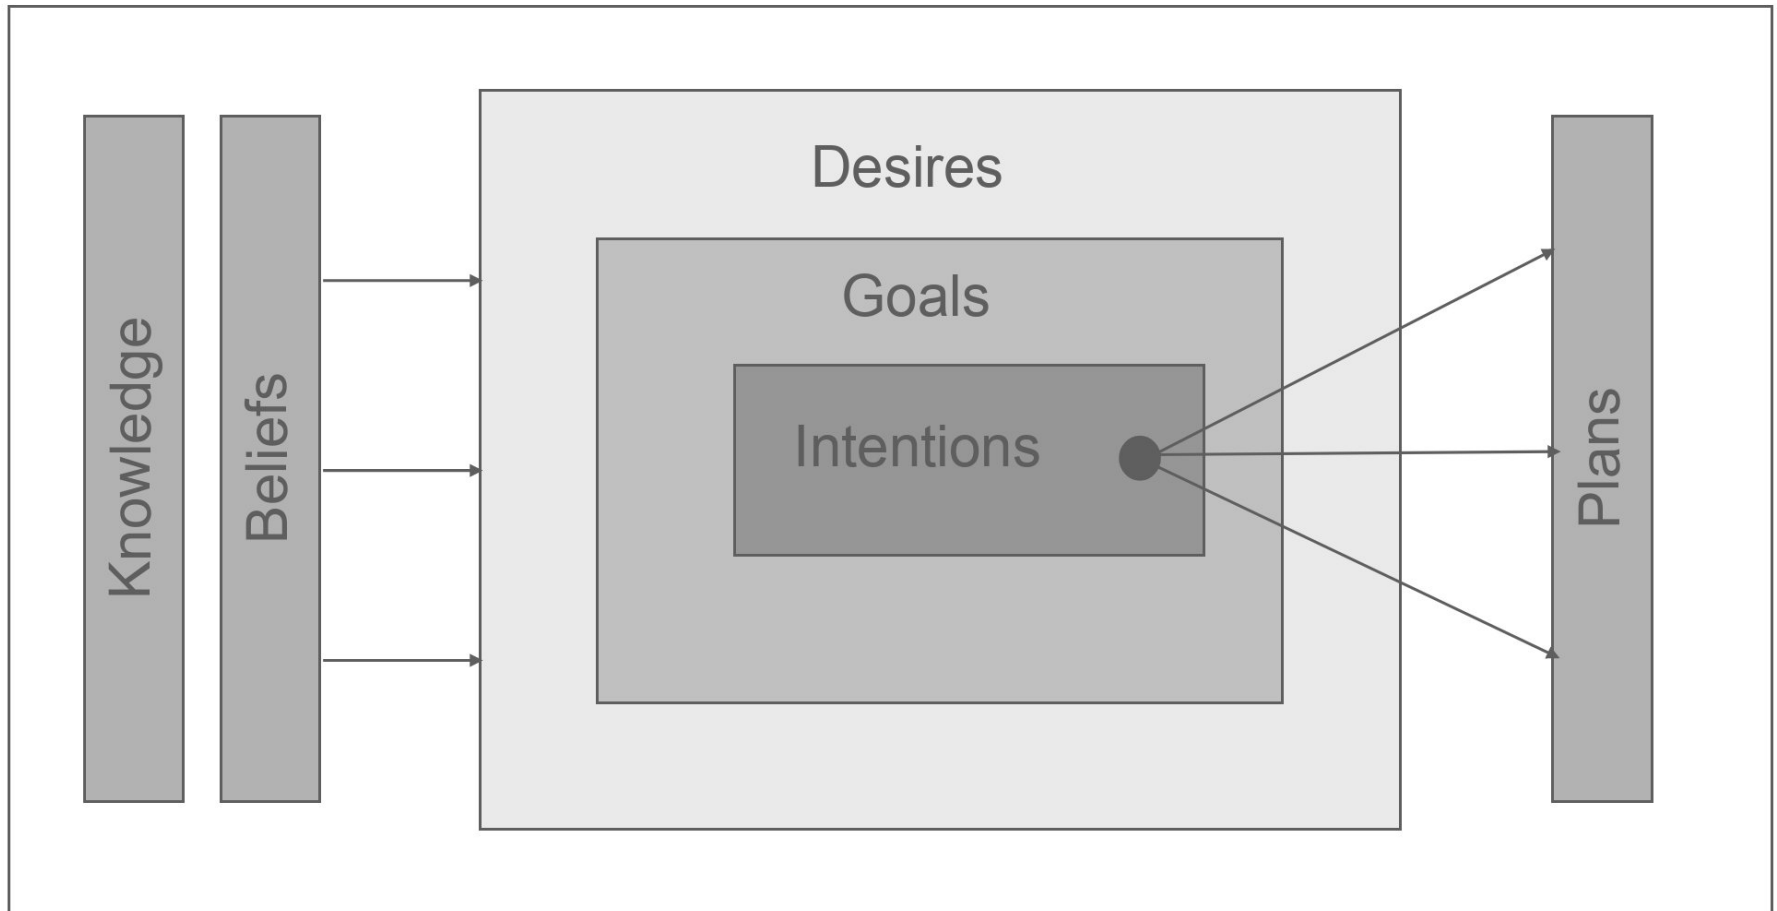
Belief-Desire-Intention paradigm



Human practical plans

- In BDI implementations plans usually have:
 - *Name*
 - *Goal*: invocation condition that is the triggering event for the plan
 - *Pre-condition list*: list of facts which must be true for the plan to be executed
 - *Delete list*: list of facts that are no longer true after the plan is performed
 - *Add list*: list of facts made true by executing the actions of the plan
 - *Body*: list of actions

Belief-Desire-Intention architecture



Intention is choice with commitment (Cohen & Levesque)

- Intentions (& plans) enable the agent to be **goal-driven** rather than event-driven
- By committing to intentions the agent can pursue **long-term goals**
- An autonomous agent should act on its intentions, not in spite of them
 - Adopt intentions that are feasible
 - Drop the ones that are not feasible
 - Keep (or commit to) intentions, but not forever
 - Discharge those intentions believed to have been satisfied
 - Alter intentions when relevant beliefs change

Using plans to constrain reasoning

- An agent's **plans** serve to frame its subsequent reasoning problems so as to constrain the amount of resources needed to solve them
 - Agents *commit* to their plans
 - Their plans tell them *what to* reason about, and *what not* to reason about
- Plans can help reasoning in different levels of abstraction

Problems in the deliberative approach

- **Dynamic world**

- Update symbolic world model
- World changes while planning is being done

*More about Planning
at PAR course*

- **Representation language**

- Expressive enough to be useful in any domain
- Limited enough to be computationally tractable

- **Classical planning → complete, optimal solutions**

- High computational cost
- Sometimes a sub-optimal low-cost fast reaction can be effective

Problems in the deliberative approach

Example: Agents at Work

Goal/Desire/Need

New Beliefs/Facts

I GOT THE IDEA FOR THIS NEW SOFTWARE WHEN I WORKED WITH PATTIE MAES ON HER "AGENTS" CONCEPT AT M.I.T.

TEMPERATURE IN WORK ENVIRONMENT 12 DEGREES BELOW OPTIMUM...



Contextual Beliefs

OF COURSE MY "AGENTS" DO A BIT MORE THAN THAT...

You must need a sweater

DOWNLOADING CREDIT CARD NUMBER... MAKING PURCHASE... THANK YOU FOR SHOPPING AT L.L. BEAN.



Running Plan/Intention

WHICH IS WHY I CALL THEM "MOMS".

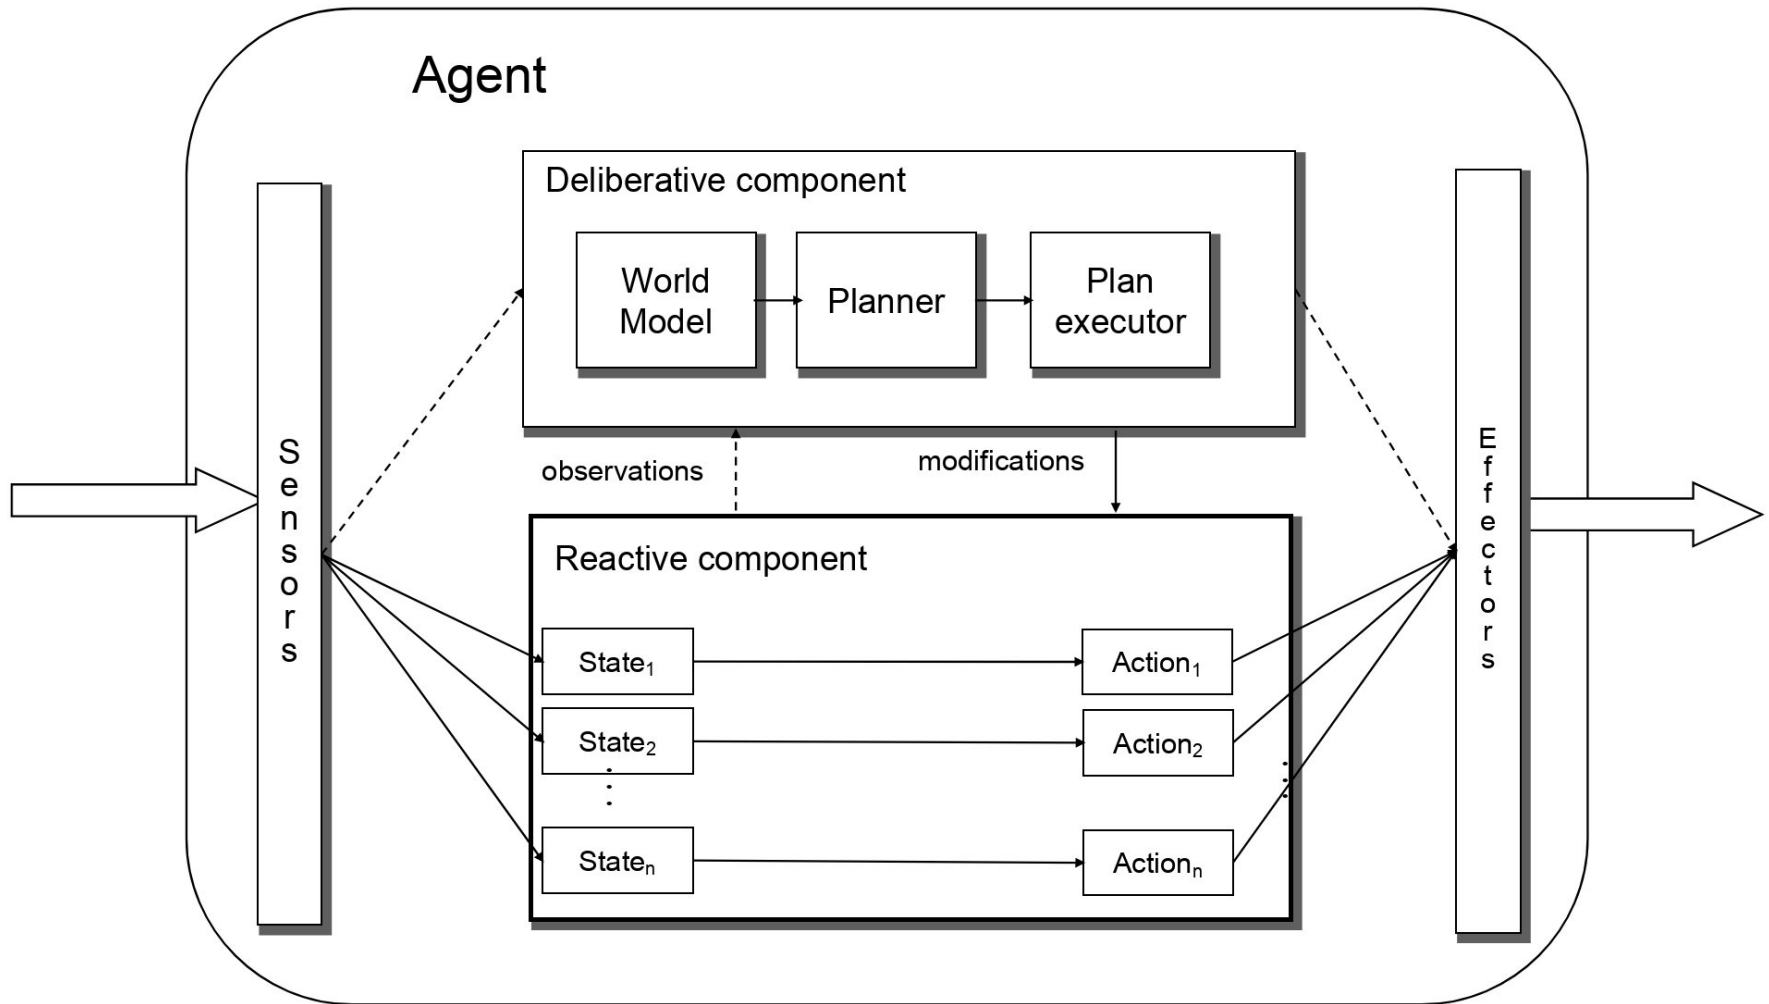
HEY, I DON'T NEED A SWEATER!



3.3. Agent architectures: Hybrid

- Many researchers have argued that neither a completely deliberative nor a completely reactive approach are suitable for building agents
- They have suggested using *hybrid* systems, which attempt to marry classical and alternative approaches
- An obvious approach is to build an agent out of (at least) two subsystems:
 - A *deliberative* one, containing a symbolic world model, which develops plans and makes decisions in the way proposed by symbolic AI
 - A *reactive* one, which is capable of reacting quickly to events without complex reasoning

Basic schema of hybrid architecture



Layered architectures


- Often, the reactive component is given some kind of precedence over the deliberative one
- This kind of structuring leads naturally to the idea of a *layered* architecture, of which TOURINGMACHINES and INTERRAP are examples
- In such an architecture, an agent's control subsystems are arranged into a hierarchy, with higher layers dealing with information at increasing levels of abstraction

Layering techniques

- A key problem in such architectures is what kind of control framework to embed the agent's subsystems in, to manage the interactions between the various layers
 - *Horizontal layering*
 - *Vertical layering*

Horizontal layering

m possible actions suggested by each layer, n layers




Each layer is directly connected to the sensory input and action output, and suggests actions to perform


Central control system can be complex, because there are $O(m^n)$ possible options to be considered

Vertical layering

m possible actions suggested by each layer, n layers



One pass



Two pass

$O(mn)$ interactions between layers

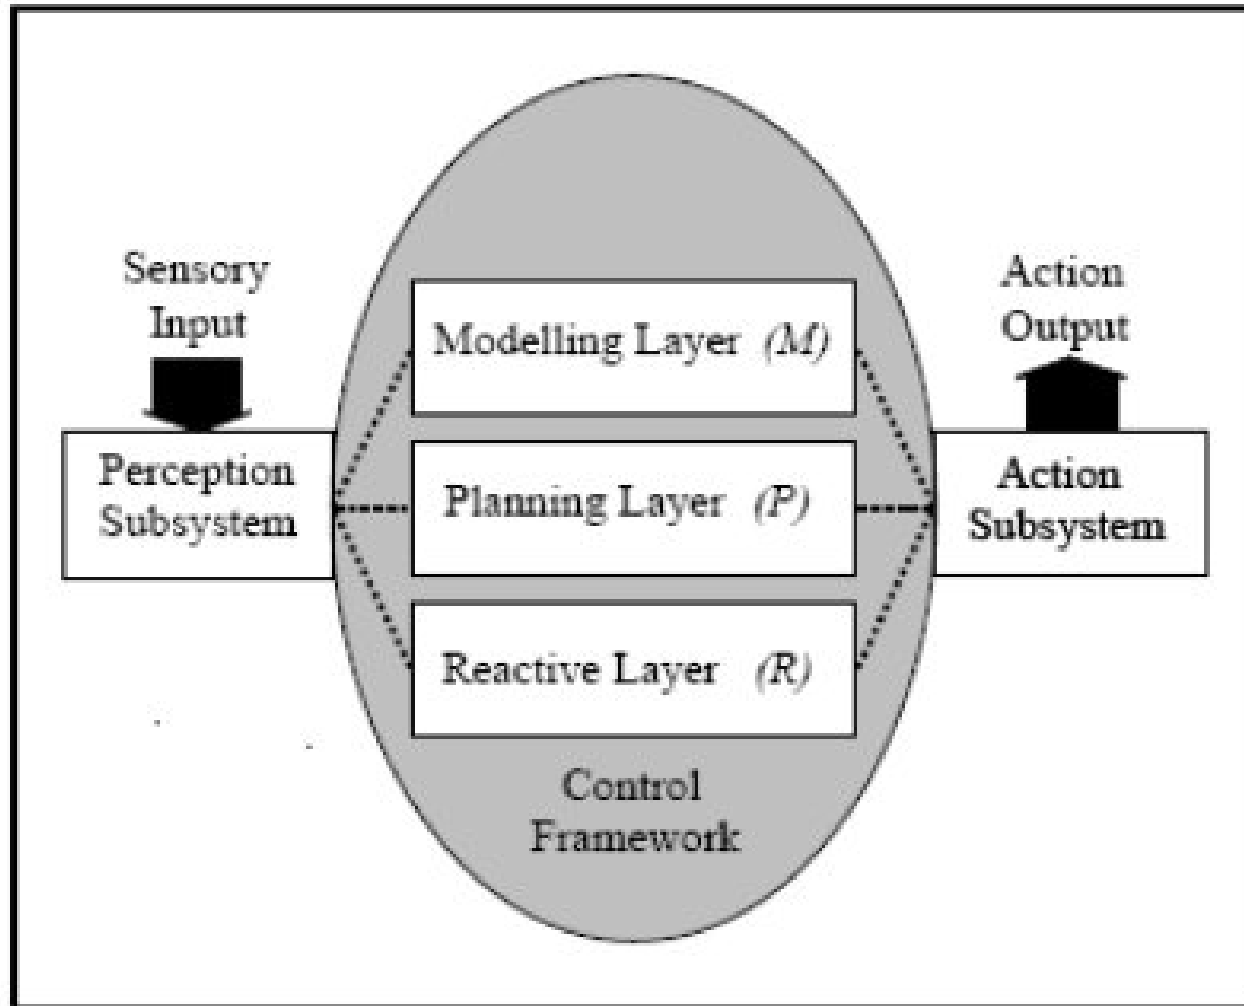
Sensory input and action output dealt with by one layer

Not fault tolerant to layer failure

Ferguson - TOURINGMACHINES

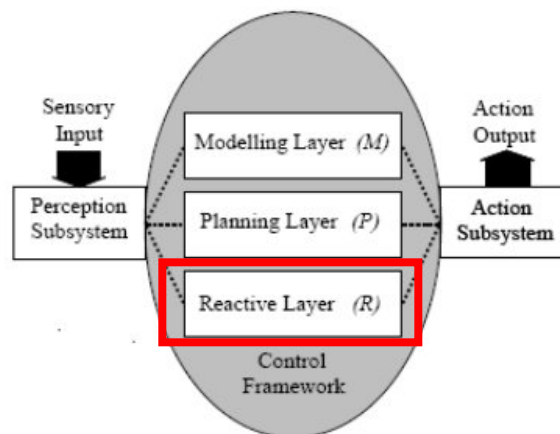
- The TOURINGMACHINES architecture consists of *perception* and *action* subsystems, which interface directly with the agent's environment, and three *control layers*, embedded in a *control framework*, which mediates between the layers

Ferguson - TOURINGMACHINES



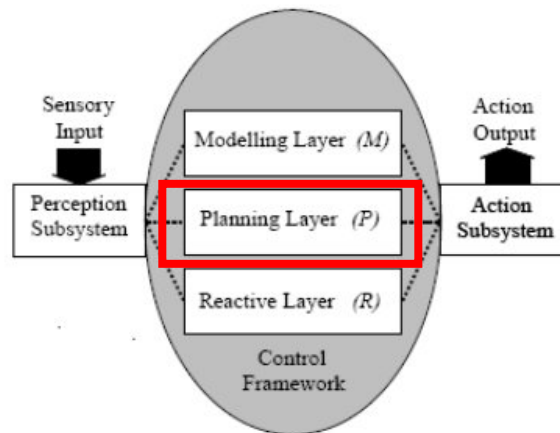
Ferguson - TOURINGMACHINES

- The *reactive layer* is implemented as a set of situation-action rules (subsumption architecture)
rule-1: obstacle-avoidance
if
 is-in-front(Obstacle, Observer) and
 speed(Observer) > 0 and
 separation(Obstacle, Observer) < Threshold
then
 change-orientation(ObstacleAvoidanceAngle)



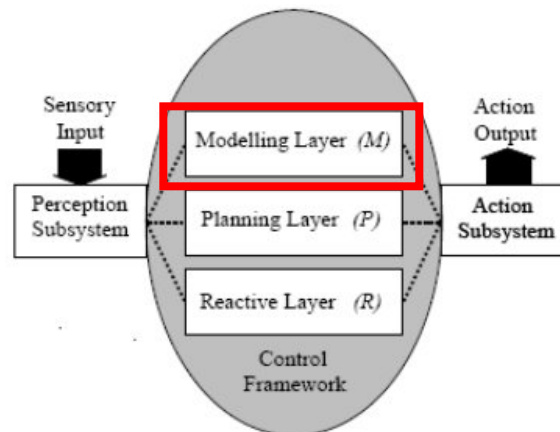
Ferguson - TOURINGMACHINES

- The *planning layer* constructs plans and selects actions to execute in order to achieve the agent's goals



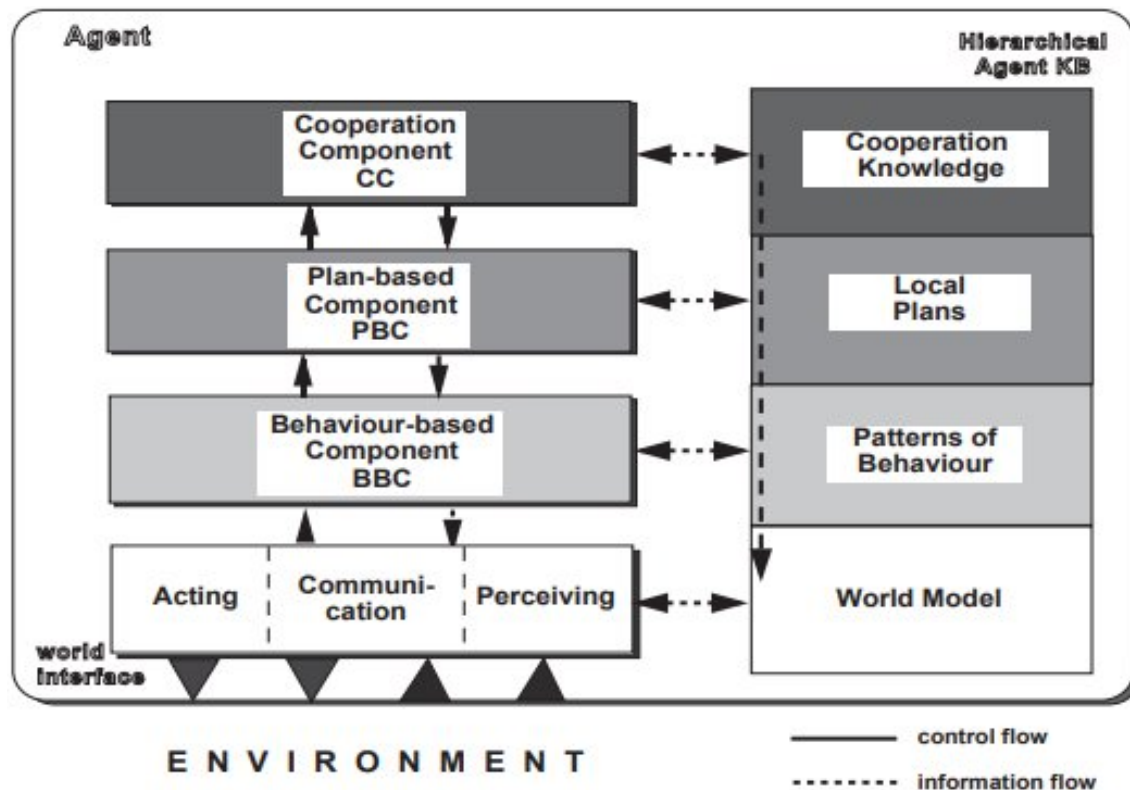
Ferguson - TOURINGMACHINES

- The *modeling* layer contains symbolic representations of the 'cognitive state' of other entities in the agent's environment
 - E.g., identifying entity behaviours or any other world events which had not been expected, detecting and resolving goal conflicts, making short- or long-term spatio-temporal predictions about entities



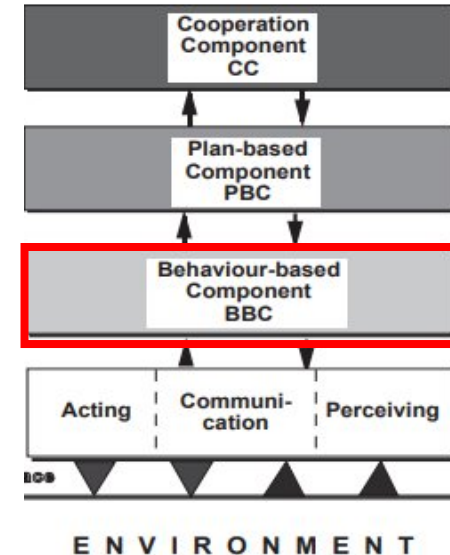
Müller - InteRRaP

- Integration of Reactive Behaviour and Rational Planning
- Vertically layered




Müller - InteRRaP: Behaviour layer (BBC)

- Reactive part of the architecture
- Works with the world model (beliefs on the world state)
- Only one level interacts with the real world
- Has a set of “situation action” rules
 - Fast recognition of situations that deserve a quick reaction
- Makes routine tasks efficiently, without complex symbolic planning



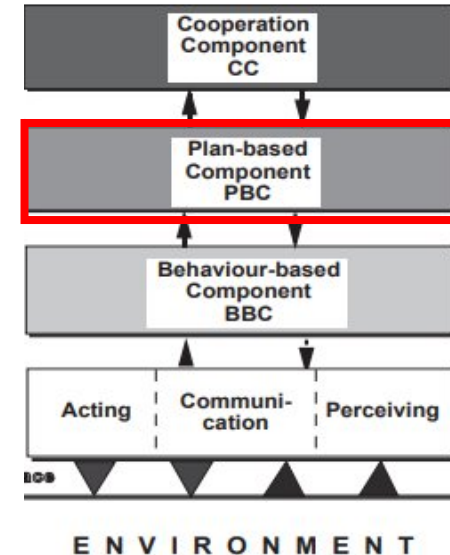
Müller - InteRRaP: Planning layer (PBC)

- Works with the mental model (beliefs on the own agent)
- Standard deliberative level
- Implements local behaviour guided towards certain goals



Müller - InteRRaP: Cooperative planning layer (CC)

- Works with the social model (beliefs on other agents of the system)
- Allows planning and cooperation with other agents
 - Global plans of action
 - Conflict resolution



Critiques to hybrid architectures

- Lack of general methodologies to guide the design process
- Very specific, application dependent
- Unsupported by formal theories

Proposed readings

- M. Wooldridge: *An introduction to Multi-Agent Systems* – *chapter 5 (reactive, hybrid)*
- A.Mas: *Agentes software y sistemas multi-agente: conceptos, arquitecturas y aplicaciones* – *chapter 2*

